

PURPOSE

Give students the minimum formal vocabulary to see social order as a *collective achievement that generates collective gain*, then use that vocabulary to frame AI alignment as the newest instance of an old problem: capturing the gains from powerful agents whose objectives we can't fully specify and whose inner states we can't inspect (“delegation under opacity”).

ENDURING UNDERSTANDING

Order among intelligent, self-interested agents is not a default — it is engineered, by three architectures (norms, markets, hierarchy) that solve the collective action problem at different scales and fail in characteristic ways. “Collective welfare” is itself a construction, and the agents are products of the very order they produce.

NARRATIVE ARC

Movement	Core question	Key concepts	Handles
1. The puzzle	Why is there order at all?	Problem of social order; collective action; free riding; coordination vs. cooperation	Hobbes, Olson, Schelling
2. Norms	How does cooperation emerge with no ruler?	Repeated games, reputation, costly signals, internalized norms, commons governance. Limit: doesn't scale past face-to-face	Axelrod, Ostrom, Durkheim
3. Markets	What coordinates strangers at scale?	Prices as information; knowledge problem; Pareto improvement. Limits: externalities, public goods	Smith, Hayek, Pareto
4. Hierarchy	Why do firms and states exist?	Transaction costs; authority & legitimacy; principal-agent problem; Goodhart's law; legibility and its costs	Coase, Weber, Scott
5. Aggregation	What does “the group” want?	Preference cycling; micromotives → macrobehavior; exit & voice; whose order is it?	Arrow, Schelling, Hirschman, Marx
6. Selves	Where do the agents come from?	Socialization, roles, position-shaped interests — groups make selves make groups	Durkheim, Marx

BRIDGE TO THE FOUR ALIGNMENT TALKS

The unit is the formal debrief of the four narrative talks (Russell/AI; the 1620 chartered corporation; the Neolithic expert; the earliest-humans “other minds”). Mapping: **corporation** → principal-agent + Goodhart (Mvt 4); **experts** → signaling, reputation, verification (Mvt 2, 5); **other minds** → cooperation without enforcement (Mvt 1–2); **Russell** → the whole stack aimed at a new agent (coda).

DESIGN NOTES

- Sequence.** Run the talks first as narrative; students identify analogies, then abstract one level up (“what do all four share?”). LYSK then names what they already half-know — concepts land as recognition, not introduction.
- Assessment.** (a) Map each talk onto unit concepts. (b) Capstone prompt: *which of these mechanisms quietly assume the agent and its overseers are matched in speed and power — and what breaks if they aren't?*
- Pitfalls.** Don't let it become a theorists' tour — concepts first, names as handles. Keep coordination vs. cooperation distinct throughout; conflating them muddles which architecture solves what.
- LYSK test.** A student passes if they can state each concept in one sentence and supply one example from outside the course.